

Homework 3: An Interpreter for Mini-OCaml (MOCaml)

Due Tuesday, October 28, at 11:30pm

Turn in your homework via the course web page as an updated version of the `interp.ml` text file that I have provided in the attached zip archive. Do not rename the file, and do not turn in (or modify) any other files.

Make sure your code can be successfully loaded into the OCaml interpreter; if not you get an automatic 0 for the homework!

Recall the CS131 Academic Honesty Policy! You must say whom you discussed the assignment with at the top of your assignment, and also what other resources you used.

For this assignment, you will implement an interpreter for a subset of OCaml in OCaml. This is a fun way to get more practice with OCaml and also to really learn how the various language constructs behave. You should obey our usual style rules and tips from past homeworks, of course (e.g., no use of any imperative features of OCaml).

Now on to the assignment! First unzip the `hw3` directory from the zip archive I've provided. The directory contains an incomplete interpreter for a subset of OCaml, which we'll call Mini-OCaml, or MOCaml. Your job is to complete the interpreter by implementing the three functions declared in the file `interp.ml`. You may also declare any helper functions that you find useful in that file as well, and you may use functions from the OCaml standard library. Your MOCaml interpreter should behave just like the regular `ocaml` interpreter for the subset of OCaml that we are handling. So if you ever want to figure out what should happen when a particular expression is evaluated, just try it in `ocaml` and find out! However, there are a few important differences:

- MOCaml is *dynamically typed*: all typechecking is performed at run time. So your interpreter will happily execute any syntactically legal program (but may signal errors during execution), while OCaml will reject some programs during static typechecking. For this reason, MOCaml can allow things that OCaml does not allow, for example lists that contain elements of different types.
- You can assume no variable name is bound more than once in a pattern (while OCaml checks for this property at compile time).
- MOCaml's binary operators `<` and `=` apply only to integers (you should raise a dynamic type error if the operands are not both integers), while they are more general in OCaml.
- The error messages displayed by your interpreter will be different than in OCaml.
- You can exit your interpreter by typing a single period followed by a return.

Here is a description of the various files I've provided and what you must do. **Read the rest of this document carefully and completely before getting started!**

- `hw3.ml`: This file simply loads all of the other files in order of dependencies. Therefore, to load your MOCaml interpreter, type the directive `#use hw3.ml;;` in the OCaml interpreter. Now type `mocaml();;` to fire up your interpreter or `runtests();;` to run your test suite on your code (see below for details). *You do not need to understand the code in this file.*
- `repl.ml`: This file contains code that implements the MOCaml *read-eval-print loop* or REPL, which continually reads a declaration, parses it, evaluates it by calling your `evalDecl` function (see below), and prints the result. The function `mocaml()` can be invoked to start the MOCaml interpreter. *You do not need to look at this file.*
- `test.ml`: Interpreters are complex and error-prone, so you will want to test yours extensively -- we'll be doing the same as part of grading. To facilitate testing, I've provided a simple test harness in this file -- read the comments to see how to use this test harness. The function `runtests()` will run all the tests for your interpreter. The list `tests` at the top of the file currently has a *very* small number of simple tests for your interpreter. You should add a lot more tests, for common cases, corner cases, error cases, etc. I encourage you to share your tests with others (e.g., on Piazza), to have everyone collaborate to build a comprehensive test suite. You may also find it useful to define a bunch of tests before starting your implementation, to act as a kind of "spec" for what your code should do.
- `lexer.ml` and `parser.ml`: The first step in an interpreter is to *parse* a given program, and I have done this for you. The parser checks for syntactic correctness and produces an *abstract syntax tree* that unambiguously represents the program. These two files contain a lexer and parser for MOCaml; the files were automatically generated from a high-level description of the MOCaml grammar that I wrote using the `ocamllex` and `ocamlyacc` parser generator tools. If you're curious about how these tools work, you can look at the files `lexer.mll` and `parser.mly`, which respectively describe the *tokens* in the MOCaml language and the *grammar* of the MOCaml language. *You do not need to understand any of these files in order to complete the assignment. Take CS132 if you want to implement your own lexer and parser!*
- `ast.ml`: This file defines the abstract syntax tree representation for MOCaml programs via a collection of datatypes. A MOCaml program is a single declaration, which is also a legal OCaml declaration, and its abstract syntax tree is represented as a value of type `modDecl`. *Read this file carefully in order to familiarize yourself with our representation of MOCaml programs, which you will need to manipulate in your code.*
- `env.ml`: We need a notion of *environment* to handle variables. This file contains a module `Env` that defines a (polymorphic) type `'a env` for environments along with associated operations. The type `'a env` maps strings (i.e., variable names) to values of type `'a`. *Read the comments in the file as well as the "module type" definition, which gives the interface to the environment module, in order to familiarize yourself with the type for environments and the associated operations.*
- `interp.ml` declares the three functions that you need to implement. Specifically:
 1. The `evalDecl` function evaluates a given declaration in the context of a given environment. I've implemented the case when the declaration is simply an expression: in that case we simply invoke `evalExpr` (see below) to evaluate the expression. Your job is to implement the cases for "let" and "let rec". Note that these constructs are more restricted syntactically than their counterparts in OCaml: MOCaml syntax for "let" is `let x = e` and its syntax for "let rec" is `let rec f x = e`, where `f` and `x` are both identifiers. The latter syntax ensures that "let rec" can only

used to declare functions. *Suggestion: Implement the simple "let" case, which is useful for testing purposes, and then move on to evalExpr (see below) before coming back to "let rec" later.*

2. The evalExpr function evaluates a given expression to a value, in the context of a given environment. I've implemented the simplest case for you (evaluating an integer constant) but you have to do the rest. Just like OCaml does, your function should signal a match failure exception (by raising the MatchFailure exception defined in interp.ml) whenever pattern matching fails (see the next item below). Additionally, your interpreter should perform *dynamic typechecking*, raising a DynamicTypeError exception whenever a type error is encountered (for example, when an integer is about to be added to a boolean).

The most interesting part of expression evaluation pertains to functions and function calls. A function evaluates to itself, except that a function value also stores the current environment. *This environment should be used whenever the function body is evaluated during a function call.*

Doing so will ensure that you implement OCaml's *static scoping* behavior properly.

Implementing recursion is tricky because it requires a function value's environment to refer to the function value itself, thereby requiring the creation of a cyclic data structure (which is not possible in the purely functional subset of OCaml). Fortunately there is an easy workaround for the kind of recursion we support: a recursive function's name should not get added to its own environment, but instead the function value should separately store its own name (this can be easily accomplished by the evalDecl function above). Then just before the function's body is evaluated during a function call, the mapping from the stored name to the function value should be added to the environment, ensuring that references to the function's name in its own body will be handled properly.

3. MOCaml includes the match expression along with several different kinds of patterns. You will therefore need to implement a function patMatch that matches a value against a pattern. I have declared this function for you and implemented the simplest case (matching against a pattern that is just an integer constant), but you have to do the rest. A successful match should produce an environment, which maps the variables defined in the pattern to their values. For example, in OCaml the result of matching the value [1;2;3] against the pattern x::xs is the environment mapping x to 1 and mapping xs to [2;3]. This environment is then used to allow the variables x and xs to be referenced in any expression where they are in scope. An unsuccessful match should raise the MatchFailure exception.